

Katana: A Userland Toolchain-Oriented Hotpatching System

James Oakley

14 November 2010

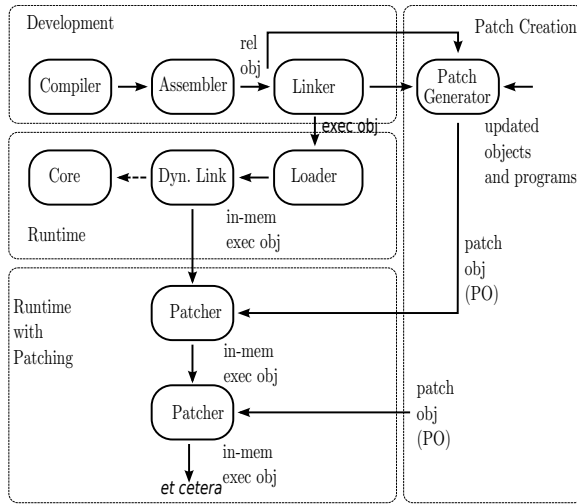
Contents

1	Introduction	1
2	Other Systems	2
3	What Katana Does	2
4	What Katana Does Not Do (Yet)	3
5	What Katana May Never Do	3
6	How to Use Katana	3
6.1	Preparing a Package for Patching Support	4
6.1.1	Source Code Practices	4
6.1.2	Compilation/Linking	4
6.2	To Generate a Patch	5
6.3	To Apply a Patch	5
6.4	To View a Patch	5
6.5	Options	5
6.6	Configuration Files	6
6.7	See Also	6
7	Patch Object Format	6
8	Patch Generation Process	6
8.1	Configuration	7
8.2	Initializing the patch object	7
8.3	Comparing source trees	7

9 Patch Application Process	7
10 Roadmap	7
11 Credits and Licensing	8

1 Introduction

Katana aims to provide a hot-patching system for userland. Further it aims to work with existing toolchains and formats so as to be easy to use and to hopefully pave the way for incorporating patching as a standard part of the toolchain. Because of this aim, Katana operates at the object level rather than requiring any access to the source code itself. This has the added bonus of making it, in theory, language agnostic (although no work has been done to test it with anything besides programs written in C). A diagram of software lifecycle with hotpatching is shown below



This document is intended to provide a users guide to Katana, insight into its inner workings, and discussion of its flaws and plans for the future. As the software is not complete, making use of Katana without understanding the inner workings and technical shortcomings is not recommended. Nevertheless, the only sections of this document necessary for “Users’ Guide” purposes are “What Katana Does”, “What Katana Does Not Do (Yet)”, and

most importantly “How to Use Katana”.

This document is a work in progress. It is not a polished guide yet.

2 Other Systems

There are other hotpatching systems in existence. The curious are invited to explore Ginseng and Polus. Both of these systems parse the source code, which adds significant complexity to them and results in significant programmer annotation of the code to give hints to the systems. Ginseng uses complicated type-wrappers when patching variables which does not fit cleanly with existing executables and has some impact on the performance of the software. Ginseng is considerably more mature than Katana, however. Neither system is production ready, but Ginseng is probably closer than Katana at the moment.

The system most like Katana in many ways is KSplice, and the curious reader is definitely invited to investigate. KSplice patches the kernel and not userland, does not attempt to patch variables, and creates patches as kernel modules rather than working towards a general ELF-based patch format.

3 What Katana Does

- Runs on x86 and x86-64
- Generates patches for simple programs
- Applies simple patches

4 What Katana Does Not Do (Yet)

- Patch any major programs: it has not yet been demonstrated on anything more than toy examples
- Provide any method to handle opaque data it cannot patch (void*, situations where which action a user would prefer is unclear, etc)
- Patch previously patched processes
- Provide robust operation
- Run on any architectures other than x86 and x86-64

- Tested on any operating system besides GNU/Linux
- Allow for calls in patched code to previously unused functions
- Work for programs which actually make use of some of the large code model features of the x86-64 ABI.
- And much more

See Roadmap for more things which are not complete

5 What Katana May Never Do

- Work on any binary formats besides ELF

6 How to Use Katana

Katana is intended to be used in two stages. The first stage generates a patch object from two different versions of an tree. By an object tree, we mean the set of object files (.o files) and the executable binary they comprise. Katana works completely at the object level, so the source code itself is not strictly required, although all objects must be compiled with debugging information. This step may be done by the software vendor. In the second stage, the patch is applied to a running process. The original source trees are not necessary during patch application, as the patch object contains all information necessary to patch the in-memory process at the object level. It is also possible to view the contents of a patch object in a human-readable way for the purposes of sanity-checking, determining what changes the patch makes, etc.

6.1 Preparing a Package for Patching Support

Katana aims to be much less invasive than other hot-patching system and require minimal work to be used with any project. It does, however, have some requirements.

6.1.1 Source Code Practices

Katana does not look at the source code, therefore unlike several other hot-patching systems, it does not require any annotation in the source code. There are, however, some best practices to follow.

- Avoid the use of `void*` at least for global variables (since Katana does not currently patch local variables, preferring to wait until any functions using changed variables are no longer on the stack). Since it is typeless and opaque, it is very hard to analyze and patch.
- Avoid unnamed types. i.e., instead of `typedef struct {...} Foo;` use `typedef struct Foo_ {...} Foo;`.
- Avoid accessing structure members by offsets instead of by the member names. As long as you keep all the code where you do this up to date, it should not be a problem, but katana cannot detect when you do this.

6.1.2 Compilation/Linking

Required CFLAGS:

- `-g`

Recommended CFLAGS:

- `-ffunction-sections`
- `-fdata-sections`

Recommended LDFLAGS:

- `-emit-relocs`

6.2 To Generate a Patch

Let the location of your project be `/project`. You must have two versions of your software available: the version identical to the running software which must be hotpatched, call it `v0`, and the version to which you wish to hotpatch the running software, call it `v1`. Let `foo` be the name of your program. Then `/project/v0/foo` must exist and `/project/v0` must also contain (possibly in subdirectories) all of the object files which contributed to `/project/v0/foo`. The source code itself is immaterial, as Katana does not parse it. Similarly, `/project/v1/foo` must exist and `/project/v1` contain all of the object files contributing to `/project/v1/foo`. Katana is then invoked as

```
katana [OPTIONS] -g [-o OUTPUT_FILE] /project/v0 /project/v1 foo
or more formally
```

```
katana [OPTIONS] -g [-o OUTPUT_FILE] OLD_OBJECTS_DIR NEW_OBJECTS_DIR
EXECUTABLE_NAME
```

If `-o OUTPUT_FILE` is not specified, the output file will be `OLD_OBJECTS_DIR/EXECUTABLE_NAME.po`

6.3 To Apply a Patch

The process to be patched is running with a pid of PID. It can be patched from its current version to a more recent version by the Patch Object (PO) file PATCH. Katana is then invoked as

```
katana [OPTIONS] -p [-s] PATCH PID
```

If all goes well, the patcher will run, print out some status messages, and leave your program in better state than it found it. The optional -s flag tells Katana to stop the target program after patching it and detaching from it. This is mostly of use for debugging Katana.

6.4 To View a Patch

One of the goals of Katana and its Patch Object (PO) format is to increase the transparency of patches: a user about to apply a patch should know what it will do. This goal is not yet fully realized, but it is possible to view some information about a patch with

```
katana [OPTIONS] -l PATCH
```

6.5 Options

The following options may be passed to katana regardless of whether one is generating, applying, or viewing a patch:

- -c CONFIG where CONFIG is the name of a configuration file to load

6.6 Configuration Files

Katana loads configuration files as follows. Configuration files loaded later in the sequence may overwrite settings from files earlier in the sequence.

- *etc/katana* + *~.katana*
- *~/.config/katana*
- *./katana*
- any file specified with -c

Configuration files are written in JSON. The JSON requirement that strings be quoted is relaxed (i.e. anything is assumed to be a string unless it can be interpreted otherwise). The following properties are recognized:

- `maxWaitForPatching <INTEGER>` This value specifies the maximum number of seconds to wait for the target to enter a safe state.
- `flags <OBJECT>` The value of flags should be an object which may contain the following properties, all of which should be bool-valued:
 - `checkPtraceWrites` Whenever something is written into the target memory, read the value back out and verify that it was written correctly. This has a performance penalty, but does provide some more robust error checking, although it should not be necessary.

6.7 See Also

the `katana` manpage (although the information in this document is considerably more extensive than in the manpage)

7 Patch Object Format

This section of the document is not yet written. It will provide a description and specification of the PO format used by Katana

8 Patch Generation Process

This section of the document is still under construction. When complete, it will provide a description of the internal process that Katana uses to generate a patch. Understanding it is not necessary for using Katana.

8.1 Configuration

Katana reads configuration files from (in order, with later configuration files overriding options found in earlier ones) from `/etc/katana`, `~/.katana`, `~/.config/katana`, and `./.katana`.

8.2 Initializing the patch object

Katana sets up a patch object ELF file with the necessary sections, see Patch Object Format

8.3 Comparing source trees

- Katana compare the old and new source trees, looking at the object (.o) files.
- For object files which exist only in the new tree, their contents are added to the patch object being created.
- For object files which exist only in the old tree, a warning about their removal is issued and nothing further is done.
- For object files which exist in both trees, type diffing and function diffing are performed and the differences are written to the patch object being created.

9 Patch Application Process

This section of the document is not yet written. It will provide a description of the internal process that Katana uses to apply a patch. Understanding it is not necessary for using Katana.

10 Roadmap

This section is highly incomplete. Future goals include

- Better interaction with the heap and dynamically allocated variables
- Better interaction with void*
- More efficient use of .rodata
- Patching already patched processes
- Patch composition
- Patch safety checking: make sure a patch actually corresponds to the process it's being applied to
- Storing warnings from generation inside a patch

11 Credits and Licensing

Katana is under development at Dartmouth College and Copyright 2010 Dartmouth College. It may be distributed under the terms of the GNU General Public License with attribution to Dartmouth College as specified in the file COPYING distributed with Katana. This document is Copyright 2010 Dartmouth College and may be distributed under the terms of the GNU Free Documentation License as found in the file FDL which should have been distributed with this documentation. If it was not, it may be found at <http://www.gnu.org/licenses/fdl.txt>.

Katana is being written by James Oakley and was designed by Sergey Bratus, Ashwin Ramaswamy, James Oakley, Michael Locasto, and Sean Smith.